

Stack

The **runtime stack** is a **last-in-first-out (LIFO)** data structure that is used to store function arguments, return values, and local variables.

A **push operation** decrements the **stack pointer (ESP)** by 4 bytes and copies a 32-bit value into the memory location pointed to by ESP.

A **pop operation** removes a 32-bit value from the memory location pointed to by ESP and increments ESP by 4 bytes.

=====

Push Operation:

=====

The following assembly code shows how to push and pop a 32-bit value onto and off of the stack:

```
; Push the value 10 onto the stack.  
push 10  
  
; Pop the value off of the stack and store it in the register EAX.  
pop eax
```

Yes, the stack grows downwards in memory. This means that when you push a value onto the stack, it is stored at a lower memory address than the previous value on the stack.

This is different from the way that a pile of books grows. When you put a book on top of a pile of books, the pile grows upwards.

When we say that a new item pushed onto the stack has a lower memory address, are we implying that the stack starts from the very high memory regions and moves to the very low memory regions?

For instance, if I have 4GB of RAM and a stack of 4096 bytes, does the stack start at the highest memory address, like 4096, and when I push a value like 10 onto the stack, the stack pointer points to 4092, and if I push 11 onto the stack, it goes to 4088?

Then, if I were to pop 10 from the stack, does the stack pointer move back to 4092? Is this the conventional behavior of a stack in memory allocation?

Yes, you're on the right track in understanding how the stack works in memory allocation. When we say that a new item pushed onto the stack has a lower memory address, it implies that the stack grows from high memory addresses to low memory addresses.

Let me break down your example with your 4GB RAM and a stack size of 4096 bytes to clarify:

1. Initially, the stack pointer points to the top of the stack, which is the highest memory address, let's say 4096 (assuming the stack starts at the highest memory address).
2. You push the value 10 onto the stack. The stack pointer moves down by 4 bytes(for the esp is 32 bits) - 4092 is the new ESP location.
3. Then, you push the value 11 onto the stack. The stack pointer(ESP), moves down further to 4096 - 4 (4088 in your example).
4. If you were to pop the value 10 from the stack, the stack pointer would move back up to 4GB - 5 (4092), effectively "removing" the top element.

So, in summary, as you push items onto the stack, the stack pointer moves to lower memory addresses. When you pop items from the stack, it moves back up.

This is a fundamental concept in stack data structures, and it ensures that the most recently added item is always at the top of the stack.

FIGURE 5–3 Pushing integers on the stack.

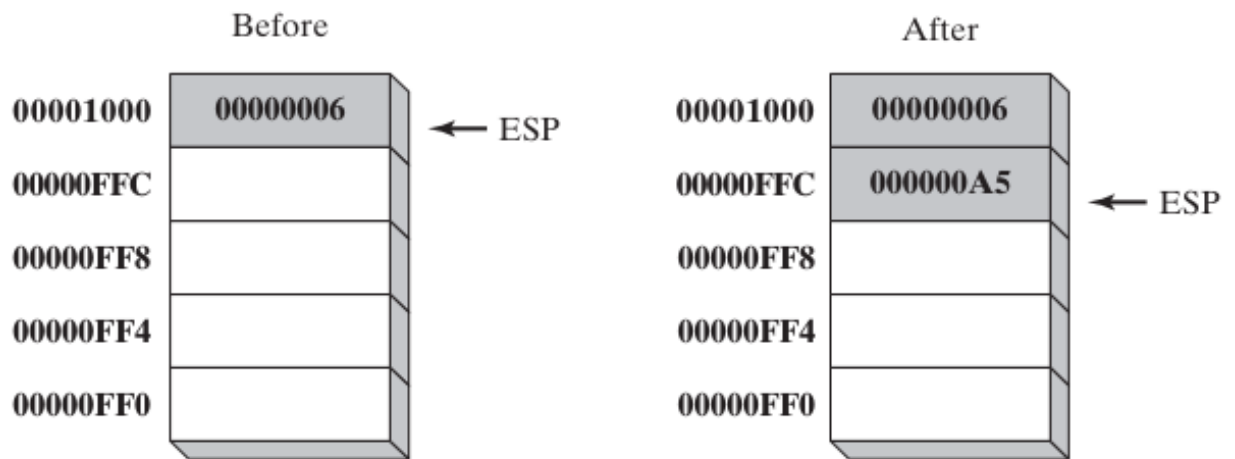



FIGURE 5–4 Stack, after pushing 00000001 and 00000002.



The push operation in a 32-bit system involves two main steps:

- 1. Decrementing the Stack Pointer:** The stack pointer (ESP) is decremented by 4 bytes.
- 2. Copying the Value:** The value you want to push is copied into the memory location pointed to by the updated stack pointer.

This effectively adds the value to the top of the stack.

The stack pointer always points to the last item pushed onto the stack.

The runtime stack typically grows downward in memory, from higher memory addresses to lower memory addresses.

This means that as you push items onto the stack, the stack pointer moves to lower memory addresses.

=====

Pop Operation:

=====

The pop operation **removes a value from the stack and increments the stack pointer.**

The pop operation in stack management is crucial for removing a value from the stack and restoring the stack pointer to its appropriate location. Here's a breakdown of how the pop operation works:

Removing a Value: When you perform a pop operation, the value at the memory location pointed to by the stack pointer is removed from the stack. This effectively takes the top item off the stack.

Incrementing the Stack Pointer: After the value is popped from the stack, the stack pointer (ESP) is incremented by the stack element size. This size is typically 4 bytes in a 32-bit system because it corresponds to the size of a 32-bit integer.

Let's use the example you provided:

Suppose you have a stack with the value 00000002 at the top, and you want to perform a pop operation to remove it.

Before the pop: ESP points to the memory location where 00000002 is stored. After the pop:

The value 00000002 is removed from the stack. ESP is incremented by 4 bytes (the size of a 32-bit integer).

ESP now points to the next-highest location in the stack.

Regarding the use of PUSHFD and POPFD to save and restore flags in assembly language, it's a useful technique.

Flags are often important for controlling program flow and making decisions based on certain conditions.

By pushing the flags onto the stack and later popping them, you can effectively save and restore the flags' state.

However, it's essential to be cautious when using such instructions to ensure that the program's execution path does not skip over the **POPFD instruction**.

Documentation and careful programming are critical in this regard.

An alternative method, as you mentioned, is to **push the flags onto the stack and immediately pop them into a variable**.

This provides a more structured and **less error-prone way** to save and restore the flags, as you can clearly see where the flags are being stored and retrieved.

In summary, the pop operation in stack management is used to remove values from the stack and increment the stack pointer accordingly.

Techniques like using PUSHFD and POPFD or pushing flags onto the stack and immediately popping them into variables can be valuable for managing program flags in assembly language.

=====

PUSHFD and POPFD Operation:

=====

```

.data
savedFlags DWORD ?

.code
main PROC
    ; Save the flags using PUSHFD
    pushfd

    ; Store the saved flags into the 'savedFlags' variable
    pop savedFlags

    ; Modify some flags (for demonstration purposes)
    cli ; Clear Interrupt Flag
    sti ; Set Interrupt Flag

    ; Restore the original flags using 'savedFlags'
    push savedFlags
    popfd

    ; Your code continues here

    ; Exit the program
    mov eax, 0x4C00 ; Exit code 0
    int 0x21        ; Call DOS function to exit
main ENDP

END main

```

The **PUSHFD instruction** in x86 assembly language knows where to find the flags in the flags register. You don't need to specify anything explicitly when using PUSHFD.

The **PUSHFD instruction** is designed to push the entire contents of the flags register (EFLAGS) onto the stack, and it does so automatically.

It saves all the status flags, such as the Zero Flag (ZF), Carry Flag (CF), Overflow Flag (OF), etc., as well as other control flags like the Direction Flag (DF) and Interrupt Flag (IF).

When you execute PUSHFD, it takes care of pushing the 32-bit value representing the flags onto the stack, and you don't need to specify individual flags.

Later, when you use POPFD, it restores the flags register from the value on the stack.

So, you can use PUSHFD and POPFD without specifying the flags explicitly; they handle the flags manipulation for you.

In this example:

We declare a DWORD variable named savedFlags in the .data section to store the saved flags.

Inside the main procedure, we first use PUSHFD to push the flags onto the stack.

Then, we immediately POP the saved flags from the stack and store them in the savedFlags variable.

Next, we modify some flags for demonstration purposes. We clear the Interrupt Flag (CLI) and then set it (STI). These changes will be temporary.

To restore the original flags, we push the savedFlags variable onto the stack and then use POPFD to copy the saved flags back into the flags register.

Your code can continue after this point.

Finally, we use the DOS interrupt int 0x21 to exit the program. The value 0x4C00 in eax specifies that the program should exit with code 0.

Please note that using PUSHFD and POPFD to manipulate flags should be done carefully, as it affects the control flow and behavior of your program. Ensure that you save and restore the flags correctly to avoid unintended consequences.

```
; Reversing a String (RevStr.asm)
.386
.model flat,stdcall
.stack 4096
ExitProcess PROTO,dwExitCode:DWORD
```

```
.data
aName BYTE "Abraham Lincoln",0
nameSize = ($ - aName) - 1
```

```
.code
main PROC
    ; Initialize registers and counters
    mov ecx,nameSize
    mov esi,0
```

```
L1:
    ; Get the next character from the 'aName' string
    movzx eax,aName[esi]

    ; Push the character onto the stack
    push eax

    ; Move to the next character
    inc esi

    ; Continue the loop until all characters are processed
    loop L1

    ; Pop the characters from the stack (in reverse order)
    ; and store them back into the 'aName' array

    ; Reset counters
    mov ecx,nameSize
    mov esi,0
```



```

L2:
    ; Get the character from the stack
    pop eax

    ; Store the character in the 'aName' string
    mov aName[esi],al

    ; Move to the next character
    inc esi

    ; Continue the loop until all characters are restored
    loop L2

    ; Exit the program
    INVOKE ExitProcess,0

main ENDP
END main

```

Now, let's break down what this code is doing:

The code begins with some assembly directives and the declaration of the `ExitProcess` function, which is used to exit the program later.

In the `.data` section, there's a string `aName` containing "Abraham Lincoln," and `nameSize` is calculated to determine the length of the string.

In the `.code` section, the main procedure starts. It initializes some registers and counters for later use.

The first loop labeled as `L1` processes each character in the `aName` string:

It retrieves a character from the string.

Pushes that character onto the stack.

Moves to the next character in the string.

Repeats until all characters are processed.

After pushing all characters onto the stack, the code proceeds to the second loop labeled as L2, which restores the characters in reverse order:

It pops a character from the stack.

Stores that character back into the aName string. Moves to the next character. Continues until all characters are restored. Finally, the program invokes the ExitProcess function to exit.

This program effectively reverses the characters in the aName string by pushing them onto the stack and then popping them off in reverse order. The stack's Last-In-First-Out (LIFO) behavior ensures that the characters are reversed in the string.

Here's what's happening step by step:

.data Section:

In the .data section, you've defined a null-terminated string named aName with the value "Abraham Lincoln" and terminated with a null character 0. This string represents the name you want to reverse. You've also calculated the length of the string using nameSize.

The expression (\$ - aName) calculates the difference between the current location (\$, which represents the address of the current instruction) and the start of the aName string. This effectively gives you the length of the string. However, you subtract 1 to account for the null terminator, as it's not part of the string length.

.code Section:

In the .code section, the main procedure begins. mov ecx, nameSize: This instruction moves the value of nameSize (which represents the length of the string) into the ecx register. The ecx register is often used as a loop counter or for holding counts in assembly language. mov esi, 0: This instruction initializes the esi register to 0. The esi register is often used as a pointer or an index when working with strings or arrays. So, in summary, this part of the code sets up the ecx register with the length of the aName string and initializes the esi register to 0, which is used to iterate through

the string character by character in subsequent parts of the code.

L1: Reversing the String and Pushing onto the Stack

In this part of the code labeled as L1, the goal is to reverse the characters in the `aName` string and push them onto the stack in reverse order:

`movzx eax, aName[esi]`: This instruction loads the next character from the `aName` string into the `eax` register. `movzx` is used to zero-extend the character from a byte to a 32-bit register. `aName[esi]` accesses the character at the current position pointed to by `esi`.

`push eax`: After loading the character into `eax`, it is pushed onto the stack using the `push` instruction. This effectively reverses the order of characters since the stack is a Last-In-First-Out (LIFO) data structure.

`inc esi`: This increments the `esi` register, which acts as an index, to move to the next character in the string.

`loop L1`: The `loop` instruction decrements the `ecx` register (loop counter) and checks if it's zero. If `ecx` is not zero, it jumps back to L1 to process the next character. This loop continues until all characters are pushed onto the stack.

L2: Popping Characters from the Stack and Restoring the String

In this part of the code labeled as L2, the reversed characters are popped from the stack and stored back into the `aName` string:

- **`pop eax`:** The `pop` instruction retrieves a character from the stack (in reverse order) and stores it in the `eax` register.
- **`mov aName[esi], al`:** This instruction takes the character in `al` (the lower byte of `eax`) and stores it back into the `aName` string at the current position pointed to by `esi`.
- **`inc esi`:** `esi` is incremented to move to the next character in the string.
- **`loop L2`:** The loop continues until all characters have been restored, just like in L1.

Finally, the program exits using `INVOKE ExitProcess, 0`.

Overall, this code snippet efficiently reverses the characters in the `aName` string by pushing them onto the stack in `L1` and then popping them off in reverse order in `L2`, effectively reversing the string in place.

Certainly, let's focus on the more complex parts of the code and break down the challenging aspects:

String Processing: The code is manipulating a string (`aName`) character by character. This involves understanding how strings are represented in memory and how to access individual characters using indexing (e.g., `aName[esi]`). The `movzx` instruction is used to zero-extend the character from a byte to a 32-bit register (`eax`) for proper handling.

Looping with ECX: The code uses the `ecx` register as a counter to control loops. Understanding how loop counters work and how they decrement with each iteration is essential. The `loop` instruction is used to check the counter and decide whether to continue looping.

Stack Manipulation: Pushing and popping values onto/from the stack can be conceptually challenging. It's important to grasp that the stack is a last-in-first-out (LIFO) data structure. `push` pushes a value onto the stack, and `pop` retrieves it. In this code, characters are pushed onto the stack in one loop and popped from it in another loop to reverse their order.

Register Usage: The code uses various registers (`eax`, `esi`, `ecx`, `al`) for different purposes. Keeping track of what each register is used for and how they interact with the data is crucial.

Procedure Calls: The program uses the `INVOKE` macro to call the `ExitProcess` function to exit. Understanding how to make function calls and pass parameters is fundamental in assembly programming.

Memory Access and Data Movement: Manipulating memory and moving data between registers and memory locations requires a solid understanding of assembly concepts. This includes addressing modes, data sizes (byte, word, dword), and operations like `mov` (move) and `movzx` (move with zero extension).

with zero extension).

String Length Calculation: The calculation of `nameSize` may be challenging for beginners. It determines the length of the string (`aName`) by subtracting the address of the string's beginning from the current location (`$`). This can be a bit abstract and requires knowledge of how memory addresses work.

Control Flow: Understanding how the program flows from one section to another, particularly the branching between loops (`L1` and `L2`), is essential.

These are some of the more challenging aspects of the code. Assembly language programming can be intricate because it requires precise management of low-level details. However, with practice and a good grasp of the fundamentals, it becomes easier to work with.

Let's review the questions one by one:

Which register (in 32-bit mode) manages the stack?

In 32-bit mode, the ESP (Extended Stack Pointer) register manages the stack. It points to the top of the stack, and as items are pushed or popped, it is adjusted accordingly.

How is the runtime stack different from the stack abstract data type?

The runtime stack, as used in computer programming, is a memory region used for function calls, local variable storage, and managing program flow. It grows and shrinks as functions are called and return. The stack abstract data type, on the other hand, is a concept in computer science that represents a data structure with a specific set of operations (push and pop) and follows the Last-In-First-Out (LIFO) principle. While the runtime stack is implemented using memory in the computer's architecture, the stack abstract data type is a higher-level concept used in algorithms and data structures.

Why is the stack called a LIFO structure?

The stack is called a LIFO (Last-In-First-Out) structure because the last item (or element) pushed onto the stack is the first one to be popped off. It operates on the principle that the most recently added item is the one that gets accessed and removed first, similar to a stack of physical objects where you would take the top item first.

When a 32-bit value is pushed on the stack, what happens to ESP?

When a 32-bit value is pushed onto the stack, the ESP (Extended Stack Pointer) register is decremented by 4 bytes (the size of a 32-bit value). This adjustment of ESP ensures that it points to the new top of the stack.

(True/False): Local variables in procedures are created on the stack.

True. In many programming languages and on most computer architectures, local variables within functions or procedures are typically allocated on the stack. This allows for efficient memory management and scoping within function calls.

(True/False): The PUSH instruction cannot have an immediate operand.

False. The PUSH instruction can indeed have an immediate operand. You can push an immediate value onto the stack using PUSH in assembly language. For example, PUSH 42 would push the value 42 onto the stack.

Let's go through the assembly language instructions and directives you've mentioned that haven't been discussed yet:

- 1. ENDP:** This directive marks the end of a procedure. It is used to define the end of a named procedure block. Typically, you see it paired with PROC to define the beginning and end of a procedure.
- 2. POP:** The POP instruction is used to remove (pop) the top value from the stack and store it in the specified destination operand. For example, POP RAX would pop the top value from the stack and store it in the RAX register.
- 3. POPA and POPAD:** These instructions are used to pop all general-purpose registers in a specific order (from DI to AX for POPA and from EDI to EAX for POPAD) from the stack.
- 4. POPFD:** This instruction pops the values of the EFLAGS register (the processor flags) from the stack. It is used to restore the state of the flags.
- 5. PROC:** The PROC directive is used to define the beginning of a named procedure block. It's often paired with ENDP to enclose the code of a subroutine or function.

6. PUSH: The PUSH instruction is used to push a value onto the stack. For example, PUSH RAX pushes the value in the RAX register onto the stack.

7. PUSHA and PUSHAD: These instructions are used to push all general-purpose registers in a specific order (from AX to DI for PUSHA and from EAX to EDI for PUSHAD) onto the stack.

8. PUSHFD: This instruction pushes the values of the EFLAGS register (the processor flags) onto the stack. It is used to save the current state of the flags.

9. RET: The RET instruction is used to return from a procedure or subroutine. It typically pops the return address from the stack and transfers control back to the calling code.

10. USES: The USES directive is used to specify which registers a procedure uses or modifies. It helps document which registers are affected by the procedure, making it clear to the programmer.

USES: The USES directive is used in the context of procedure or subroutine declarations. It specifies which registers a procedure uses or modifies. For example, if a procedure uses the RAX and RBX registers, you can declare it like this:

```
MyProcedure PROC USES RAX, RBX
```

This helps document which registers the procedure interacts with, making it clear to other programmers.

POPA and POPAD: These instructions are used to pop all general-purpose registers in a specific order from the stack. POPA is for 16-bit registers (DI through AX), while POPAD is for 32-bit registers (EDI through EAX). They are the opposite of PUSHA and PUSHAD, which push all these registers onto the stack.

PUSHFD and PUSHAD: These instructions are used to save the state of processor flags and general-purpose registers on the stack.

PUSHFD pushes the EFLAGS register (processor flags) onto the stack. It's used to save the current state of flags. PUSHAD pushes all general-purpose registers (EAX, ECX, EDX, EBX, ESP, EBP, ESI, and EDI) onto the stack. It's used to save the values of these registers. POPFD: This instruction is used to pop the values of the EFLAGS register (processor flags) from the stack. It's used to restore the state of flags.

```
477 .model flat, stdcall
478 .stack 4096
479
480 .data
481     message db "Hello, World!", 0
482
483 .code
484     MyProcedure PROC USES EAX, EBX
485         ; Your code here that uses EAX and EBX registers
486         mov eax, 42
487         mov ebx, 24
488         add eax, ebx
489         ; Rest of your code
490         ret
491     MyProcedure ENDP
492
493     main PROC
494         ; Push all general-purpose registers onto the stack
495         pushad
496         ; Call MyProcedure
497         call MyProcedure
498         ; Pop all general-purpose registers from the stack
499         popad
500         ; Push processor flags onto the stack
501         pushfd
502         ; Modify some flags (for demonstration purposes)
503         cld ; Clear direction flag
504         ; Pop processor flags from the stack
505         popfd
506         ; Display a message
```



```

507     mov eax, 4           ; syscall number for sys_write (Linux)
508     mov ebx, 1           ; file descriptor 1 (stdout)
509     mov ecx, offset message ; pointer to the message
510     mov edx, 13          ; message length
511     int 0x80             ; invoke syscall
512     ; Exit program
513     mov eax, 1           ; syscall number for sys_exit (Linux)
514     int 0x80             ; invoke syscall
515 main ENDP
516 END main

```