

Measuring Execution Times

The code example you provided shows how to use the `GetMseconds` procedure in the `Irvine32` library to measure the execution time of a program.

The **GetMseconds** procedure returns the number of system milliseconds that have elapsed since midnight.

To measure the execution time of a program, you would first call the `GetMseconds` procedure to record the start time.

Then, you would call the program whose execution time you wish to measure. Finally, you would call the `GetMseconds` procedure again to record the end time.

The difference between the end time and the start time is the execution time of the program.

The following code example shows how to use the `GetMseconds` procedure to measure the execution time of a simple program:

```
585 .data
586     startTime DWORD ?
587     procTime  DWORD ?
588
589 .code
590     call GetMseconds
591     ; get start time
592     mov startTime, eax
593
594     ; call the program whose execution time you wish to measure
595     ; ...
596
597     call GetMseconds
598     ; get end time
599     sub eax, startTime
600     ; calculate the elapsed time
601     mov procTime, eax
602     ; save the elapsed time
```

The **variable procTime** will now contain the execution time of the program, in milliseconds.

You can use this technique to measure the execution time of any program, regardless of its complexity.

However, it is important to note that the overhead of calling the GetMseconds procedure twice is insignificant when compared to the execution time of most programs.

Relative Performance

You can also use the GetMseconds procedure to measure the relative performance of two different code implementations.

To do this, you would measure the execution time of each implementation and then divide the execution time of the first implementation by the execution time of the second implementation.

The result will be a number that indicates the relative performance of the two implementations.

For example, the following code example shows how to measure the relative performance of two different sorting algorithms:

```

627 .data
628     startTime1 DWORD ?
629     procTime1  DWORD ?
630     startTime2 DWORD ?
631     procTime2  DWORD ?
632 .code
633     ;measure the execution time of the first sorting algorithm
634     call GetMseconds    ;get start time
635     mov startTime1, eax
636     ; call the first sorting algorithm ...
637     call GetMseconds
638     ; get end time
639     sub eax, startTime1
640     ; calculate the elapsed time
641     mov procTime1, eax
642     ; save the elapsed time
643     ; measure the execution time of the second sorting algorithm
644     call GetMseconds
645     ; get start time
646     mov startTime2, eax
647     ; call the second sorting algorithm
648     ; ...
649     call GetMseconds
650     ; get end time
651     sub eax, startTime2
652     ; calculate the elapsed time
653     mov procTime2, eax
654     ; save the elapsed time
655     ; calculate the relative performance of the two sorting algorithms
656     div procTime1, procTime2
657     ; the result is now in EAX

```

The EAX register will now contain the relative performance of the two sorting algorithms. A value of 1.0 indicates that the two sorting algorithms have the same performance.

A value greater than 1.0 indicates that the first sorting algorithm is faster than the second sorting algorithm. A value less than 1.0 indicates that the first sorting algorithm is slower than the second sorting algorithm.

You can use this technique to measure the relative performance of any two code implementations, regardless of their complexity.

Comparing MUL and IMUL to Bit Shifting in Depth

In older x86 processors, there was a significant difference in performance between multiplication by bit shifting and multiplication using the MUL and IMUL instructions.

However, in recent processors, Intel has managed to greatly optimize the MUL and IMUL instructions, so that they now have the same performance as bit shifting for multiplication by powers of two.

The following code shows two procedures for multiplying a number by 36 using bit shifting and the MUL instruction:

```

661 ;Multiplies EAX by 36 using SHL, LOOP_COUNT times.
662 mult_by_shifting PROC
663 mov ecx, LOOP_COUNT
664 L1: push eax
665 ; save original EAX
666 mov ebx, eax
667 shl eax, 5
668 shl ebx, 2
669 add eax, ebx
670 pop eax
671 ; restore EAX
672 loop L1
673 ret
674 mult_by_shifting ENDP
675
676 ; Multiplies EAX by 36 using MUL, LOOP_COUNT times.
677 mult_by_MUL PROC
678 mov ecx, LOOP_COUNT
679 L1:
680 push eax
681 ; save original EAX
682 mov ebx, 36
683 mul ebx
684 pop eax
685 ; restore EAX
686 loop L1
687 ret
688 mult_by_MUL ENDP

```

The following code calls the mult_by_shifting procedure and displays the timing results:

```

692 .data
693     LOOP_COUNT = 0FFFFFFFFh
694 .data
695     intval DWORD 5
696 .code
697     call
698     GetMseconds
699     ; get start time
700     mov
701     startTime, eax
702     mov
703     eax, intval
704     ; multiply now
705     call
706     mult_by_shifting
707     call
708     GetMseconds
709     ; get stop time
710     sub
711     eax, startTime
712     call WriteDec
713     ; display elapsed time

```

The code above, is a simple example of how to measure the execution time of a program using the GetMseconds procedure in the Irvine32 library. The program multiplies the integer 5 by 36 using the mult_by_shifting procedure, and then displays the execution time.

The two .data segments in the program are used to define two variables: LOOP_COUNT and intval. LOOP_COUNT is a constant that specifies the number of times to repeat the multiplication operation. intval is the integer that is multiplied by 36.

The reason for having two .data segments is not entirely clear. It is possible that the original author of the code was simply trying to organize the data in a logical way.

However, it is also possible that the author was trying to take

advantage of some optimization in the Irvine32 library.

Regardless of the reason, it is not necessary to have two .data segments in this program. The two variables could be defined in the same .data segment without any problems.

Here is a revised version of the program with the two .data segments combined into one:

```
717 .data
718     LOOP_COUNT = 0FFFFFFFFh
719     intval DWORD 5
720 .code
721     call
722     GetMseconds
723     ; get start time
724     mov
725     startTime, eax
726     mov
727     eax, intval
728     ; multiply now
729     call
730     mult_by_shifting
731     call
732     GetMseconds
733     ; get stop time
734     sub
735     eax, startTime
736     call WriteDec
737     ; display elapsed time
```

This revised version of the program works just as well as the original version, and it is more concise and easier to read.

- You can have as many segments for .data, .code, .bss/text.
- Use segments wisely, grouping related data and code.
- Avoid excessive segments for clarity and performance.

On a legacy 4-GHz Pentium 4 processor, the `mult_by_shifting` procedure executed in 6.078 seconds, while the `mult_by_MUL` procedure executed in 20.718 seconds.

This means that using the `MUL` instruction was 241 percent slower. However, when running the same program on a more recent processor, the timings of both function calls were exactly the same.

This example shows that Intel has managed to greatly optimize the `MUL` and `IMUL` instructions in recent processors.

Therefore, there is no longer any need to use bit shifting for multiplication by powers of two.

In fact, using the `MUL` and `IMUL` instructions is generally preferred, as they are more readable and easier to maintain.